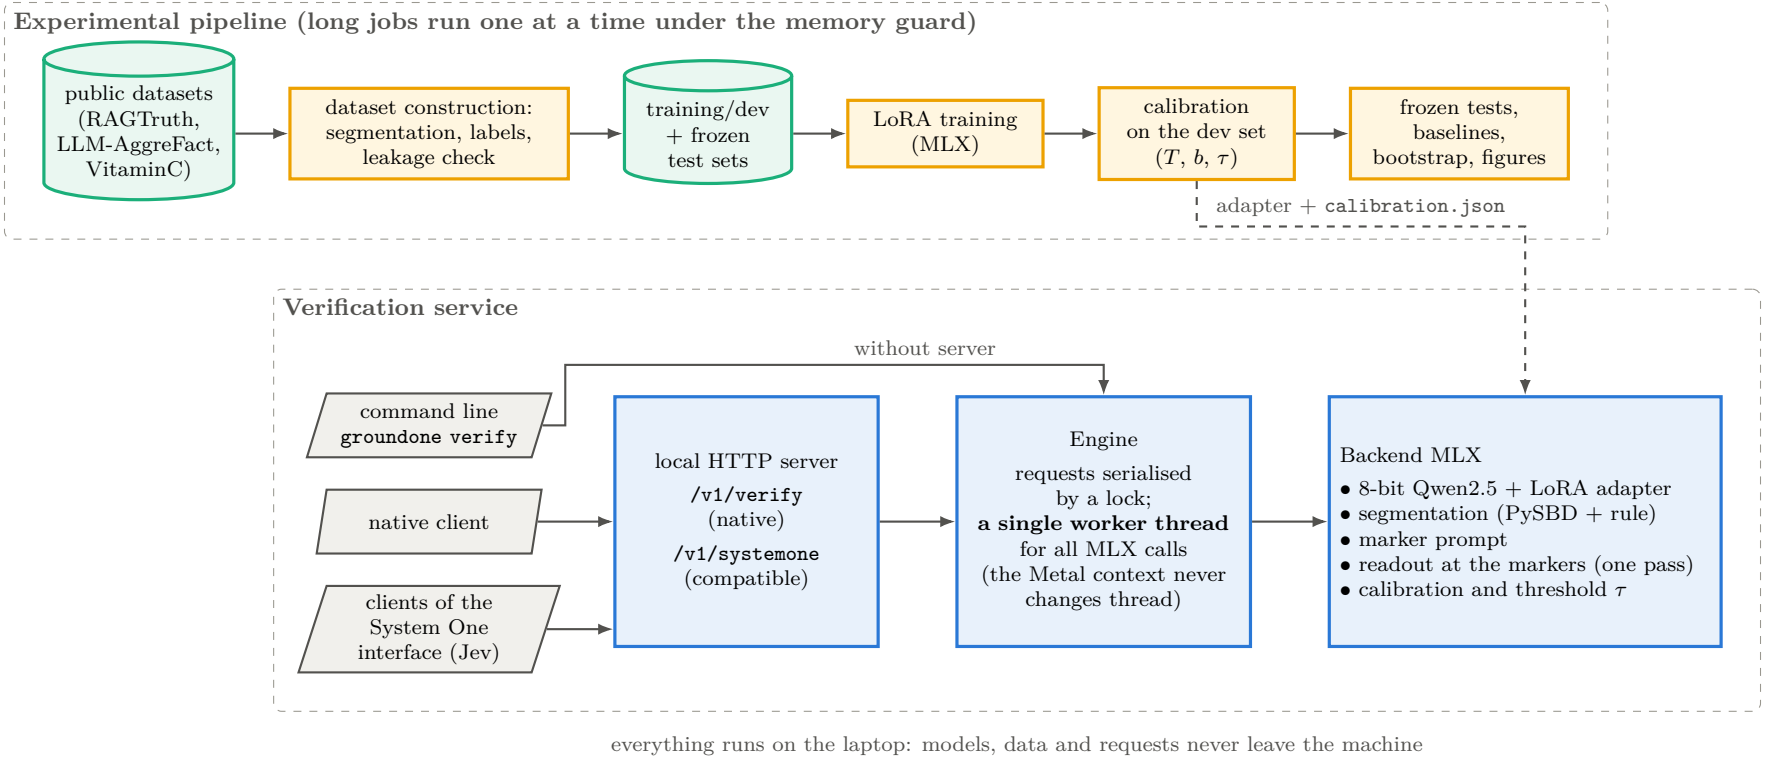

Experimental pipeline (long jobs run one at a time under the memory guard)
public datasets (RAGTruth, LLM-AggreFact, VitaminC)
dataset construction: segmentation, labels, leakage check
training/dev + frozen test sets
LoRA training (MLX)
calibration on the dev set (T, b, τ)
frozen tests, baselines, bootstrap, figures
adapter + calibration.json
Verification service
without server
command line groundone verify
native client
clients of the System One interface (Jev)
local HTTP server
/v1/verify (native)
/v1/systemone (compatible)
Engine
requests serialised by a lock;
a single worker thread for all MLX calls
(the Metal context never changes thread)
Backend MLX
• 8-bit Qwen2.5 + LoRA adapter
• segmentation (PySBD + rule)
• marker prompt
• readout at the markers (one pass)
• calibration and threshold τ
everything runs on the laptop: models, data and requests never leave the machine